

Sample Solution (Selected Questions only)

a) 1st complication: Arbitrating contention between hits & misses (or misses & misses) needs to be done. This happens because in a non-blocking cache, hits can collide with misses returning from the next level of the memory hierarchy, or misses can similarly collide. The situation is usually handled by giving priority to hits over misses, & by ordering colliding misses. ^{returning} ^{transferred}

2nd complication: to track multiple misses, so the result can be ^{transferred} to the correct cache block, & the correct pending load/store ^{can} be notified.

b) In a Write Back cache, when a dirty cache block is replaced during a Read Miss, the dirty block does not necessarily immediately move to main memory - they may wait for a long time in the Write Buffer. Hence, later again if this cache block needs to be brought back to the cache, it is not clear whether the last updates have synced to memory or not - possibly the dirty (last updated) cache block is still waiting in the Write Buffer.

c) Amdahl's law assumes that the problem size remains the same, even with the availability of multiple processors/cores. Hence, the execution time for the non-parallelizable portion of the code remains unchanged, in spite of the presence of additional hardware resources. This is not the case for Gustafson's Model, which assumes larger amount of HW resources \Rightarrow larger problem sizes, which makes the contribution of the non-parallelizable portion of the code relatively insignificant, overall leading to more optimistic estimates of speedup.

d) Advantage: if the way predictor is correct, the cache access latency is small since the hit time is small (the MUX select signal is already set to the value predicted by the predictor).

Disadvantage: There is (ideally) an 1 clock cycle penalty if the predictor is incorrect. However, in very fast processors, since the clock cycle time is extremely small, it is challenging to restrict the misprediction penalty to 1 cycle.

e) Advantage: If the branch condition is evaluated in ID (e.g. in MIPS) then an ALU operation to a register followed by a conditional branch based on this register value is a data hazard, e.g. consider the following pseudocode:

```
Sub R1, R2, R3
beq R1, zero, Label
...
```

The evaluation of the branch condition in ID is not possible, since in the same cycle, the "Sub" instruction is evaluating the R1 register in its EX stage! Hence, there must be a 1 cycle stall! (assuming EX \rightarrow ID forwarding present)

Label:

2) Let P_1 contain 1 instruction
 a) The run-time of P_1 on M_1 :

$$T_1 = \frac{(0.4 \times 2 + 0.5 \times 1 + 0.1 \times 3)I}{f_1} = \frac{(0.8 + 0.5 + 0.3)I}{2 \times 10^9} = \frac{1.6I}{2 \times 10^9} \text{ s, since } f_1 = 2 \text{ GHz.}$$

Let the frequency of M_2 be f_2 . Then, the time required to execute P_1 on M_2 :

$$T_2 = \frac{(0.4 \times 2 + 0.5 \times 2 + 0.1 \times 2)I}{f_2} = \frac{2.0I}{f_2} \text{ s}$$

Given: $T_1 = T_2 \Rightarrow \frac{1.6I}{2 \times 10^9} = \frac{2.0I}{f_2} \Rightarrow f_2 = \frac{4.0}{1.6} \text{ GHz} = 2.5 \text{ GHz.}$

b) Suppose P_1, P_2 each have "I" instructions
 M_1, M_2 each run at a frequency "f"

For M_1 :

Execution Time = $\frac{[(0.4 \times 2) + (0.5 \times 1) + (0.1 \times 3)]I}{f}$

$$T_1 = \frac{[(0.4 \times 2 + 0.5 \times 1 + 0.1 \times 3) + (0.5 \times 2 + 0.2 \times 1 + 0.3 \times 3)]I}{f} = \frac{3.7I}{f}$$

For M_2 , similarly,

$$T_2 = \frac{[(0.4 \times 2 + 0.5 \times 2 + 0.1 \times 2) + (0.5 \times 2 + 0.2 \times 2 + 0.3 \times 2)]I}{f} = \frac{4.0I}{f}$$

$\therefore M_1$ is faster overall for this workload.

c) Suppose, P_1 runs x times & P_2 runs y times

Under same assumptions at part-(b):

For M_1 : $T_1 = \frac{(1.6x + 2.1y)I}{f}$ — (1)

For M_2 : $T_2 = \frac{(2.0x + 2.0y)I}{f}$ — (2)

Given: $T_1 = T_2 \Rightarrow 1.6x + 2.1y = 2.0x + 2.0y$

$$\Rightarrow 0.1y = 0.4x$$

$$\Rightarrow x:y = 1:4$$

$\therefore P_1$ should run 1 time
 P_2 " " 4 times } Example workload.

3) a) No. of cache blocks = $\frac{\text{Cache size}}{\text{Block size}} = \frac{4KB}{16 \text{ bytes}} = \frac{2^{14} \text{ bytes}}{2^4 \text{ bytes}} = 2^8 = 256$

b) Size of array A = Size of array B = $(1024 * 4) / 16 \text{ Cache blocks} = 256 \text{ cache blocks}$
(Since each integer = 4 bytes)

Each cache block can hold: $\frac{16 \text{ bytes}}{4 \text{ bytes}} = 4 \text{ integer (array elements)}$

Starting address of A[0]: $0x00004000 = 16384$, maps to cache block: $(16384/16) \pmod{256} = 0$
~~write~~ " of A[1020]: $16384 + 1020 * 4 = 20464$, maps to cache block: $(20464/16) \pmod{256} = 255$

Starting address of B[0]: $0x00010800 = 67584$, maps to cache block: $(67584/16) \pmod{256} = 128$

" of B[1020]: $67584 + 1020 * 4 = 71664$, maps to cache block: $(71664/16) \pmod{256} = 127$

Index(i)	A access pattern	A goes to (Block #)	B access pattern	B goes to (Block #)	Comment
1023:1020	A[1023:1020]	255	B[1023:1020]	127	2 Write Misses
1019:1016	A[1019:1016]	254	B[1019:1016]	126	do
					do
515:512	A[515:512]	128	B[515:512]	0	2 write misses, B starts overwriting
511:508	A[511:508]	127	B[511:508]	255	A blocks, A starts overwriting B's blocks
					do
127:124	A[127:124]	31	B[127:124]	159	do
					do
3:0	A[3:0]	0	B[3:0]	128	

Total: (2 Write Misses) for every 4 array elements
 => $(\frac{1024}{4} * 2) = 512$ Write misses, zero read misses.

(c) # of bytes written back to memory: (write back happens when blocks start getting replaced) = $\frac{1}{2}$ of A-array + $\frac{1}{2}$ of B-array
 = $(1024 * 4) \text{ bytes} = 4 \text{ KB}$

4)

64-bit V.A.

Page size = 2kB \Rightarrow Page/frame offset = 11 bits

\hookrightarrow 2¹¹ bytes

of P.T. levels = 3.

PTE size: 8 bytes (for every level)

1st level: 4 pages

Amount of physical memory: 32GB = 2³⁵ bytes.

2nd level: 8 pages

Mask Memory Requirement: 16TB

(a) # of bits used for page offset: 11 bits.

$(64 - 11) = 2^{53}$

(b) # virtual pages/physical frames = 2

physical frames = $\frac{\text{Physical mem. size}}{\text{Frame size}} = \frac{2^{35} \text{ bytes}}{2^{11} \text{ bytes}} = 2^{24}$


(c) 1st level P.T. size: 4 pages \times (2kB/page) = 8kB = 2¹³ bytes

\therefore # of entries in 1st level P.T. = $\frac{8 \text{ kB}}{8 \text{ B}} = 2^{10}$

2nd level P.T. size = 8 pages \times (2kB/page) = 16kB

\therefore # of entries in 2nd level P.T. = $\frac{16 \text{ kB}}{8 \text{ B}} = 2^{11}$

\therefore V.A. is organized as:




$$\begin{array}{r} 2^5 \\ 2^4 + 2^{15} + 2^{25} \end{array}$$

\therefore 1st-level # of index bits: 10

2nd - " " " " : 11

3rd - " " " " : 32

(d) ~~Use of virtual address space:~~ Virtual Memory requirement: 16TB = 2⁴⁴ bytes : 2⁴⁴ bytes
 \therefore Only 44 - (11 + 10 + 11) = 12 bits of the 3rd level are used.




3rd level P.T. requires 2¹² entries

\Rightarrow Size of 3rd level P.T. = 2¹² \times (PTE size) = 2¹² \times 2³ bytes = 2¹⁵ bytes

\therefore # of pages required for 3rd level P.T. = $\frac{2^{15} \text{ bytes}}{2^{11} \text{ bytes}} = 16$

\therefore # of pages required (for 3-level P.T. Scheme) = (2⁴ + 2¹⁵ + 2²⁵) pages = 33,587,216 (pages)

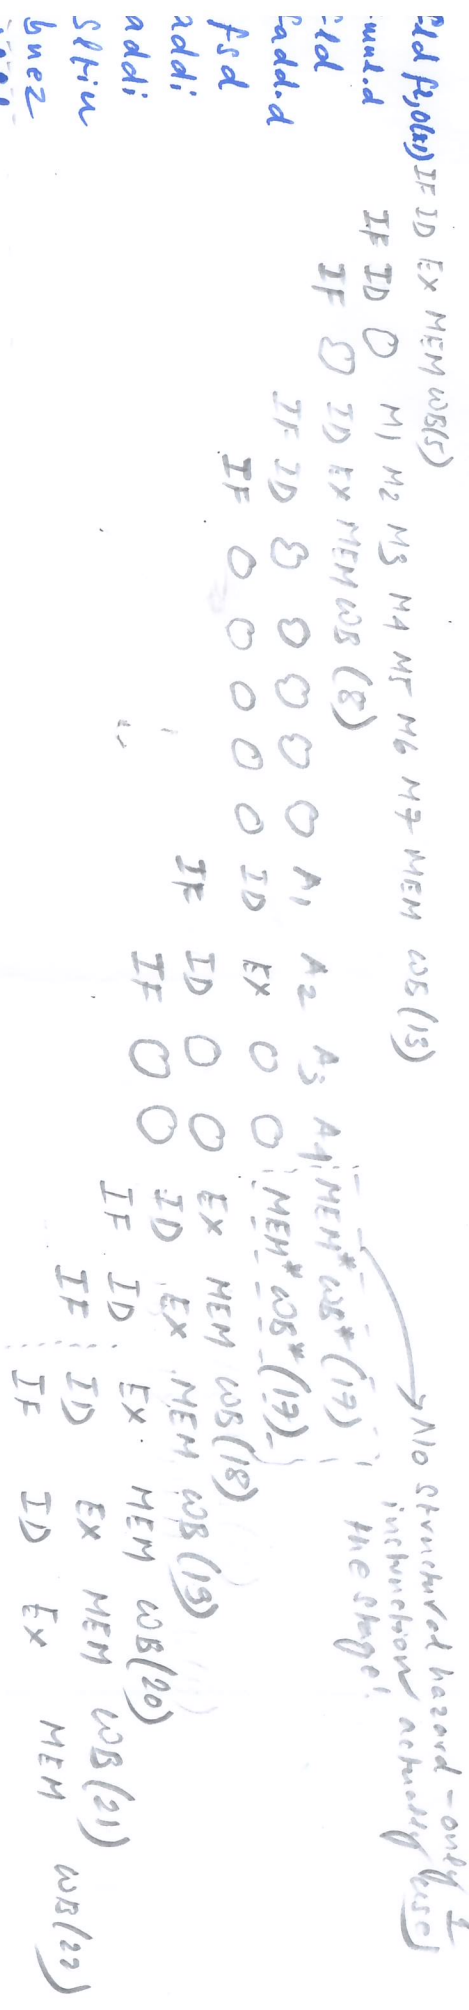
= 16 + (2¹² \times 8) + (2¹² \times 2¹¹ \times 4)



$$\frac{\text{Actual}}{\text{Flat}} = \frac{33,587,216}{2^{25}} = 0.39\%$$

(e) Flat page table Scheme would require: 2^{(44-11)}} = 2³³ entries which occupy 2³³ \times 8B frames
 (for task mentioned) [not general case] for task mentioned.

1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24 25



Notes: 1) fadd.d has latency = 6 cycles \Rightarrow takes 7 cycles to complete. EX
2) fadd.d has latency = 3 cycles \Rightarrow " 4 " " " EX

(b) # of cycles to complete a single iteration: 17 (from clock cycle #18, overlap of next iteration start)

(c) Reordered Code [with Register Renaming]

fadd.f6, f0, f8 # no stall
fadd.f6, f0, f8 # no stall, f6 renamed to f8

fadd.f6, f4, f8, f0 # no stall
addi x1, x1, 8 # no stall

addi x2, x2, 8 # no stall

fadd.f6, f4, f8 # 4 cycle stall, finishes on cycle #17
addi x3, x1, done # no stall, 4 cycle stall, finishes on cycle #15

fadd.f6, f0, f8, -8(x2), f6, # offset modified to -8, 1 cycle stall, finishes on cycle #17
fadd.f6, f0, f8, x3, f0 # finishes on cycle #18 after stalling from cycle #13

fadd.f6, f0, f8, x3, f0 # finishes on cycle #19 after stalling from cycle #14
fadd.f6, f0, f8, x3, f0 # finishes on cycle #19 after stalling from cycle #14

(d) # of cycles/iteration of re-ordered code: 13 cycles/iteration.

The above situation cannot arise in RISC-V, if EX \rightarrow EX forwarding is present.
Disadvantage: there is a 2 cycle penalty if there is a branch misprediction, in comparison to an 1-cycle penalty in MIPS.

6. Steps:

- 1) Issue
- 2) Read Operands
- 3) Execute
- 4) Write Result

See Appendix-C (C.7) of your textbook for details of these steps. You are expected to describe each of these briefly in brief.

RAW Hazards: These are avoided by making the dependent instruction wait in its functional unit, till the operand value is ready to be read from the register file (the scoreboard controller gives permission to the functional unit when it can read its value from the register file).

WAW Hazards: Before issue, the destination register for an instruction is examined. If it is the same as the destination register of a currently active instruction in any of the functional units, instruction issue stalls, & remains stalled until the active unit has completed its WB (Write Results) step.

WAR Hazards: After an instruction has completed it is checked whether an preceding instruction is yet to complete its "Read Operands" phase, with one of its operands same as the destination register of the instruction under question. If this is the case, the "Write Results" step of the instruction that has completed execution is not allowed to perform its "Write Results" — it waits till the other instruction on which it has WAR has completed its "Read Operands" stage.